

INTERNAL · OPERATING PLAN · FACELESS GENERATIVE SHORTS

Bureau of Reality: a fully generative Shorts channel, automated end to end through MCP

Version 2.2 · **Date** 06-Oct-2026 · **Owner** Sahil Mathew · **Supersedes** v2.1 (v1.0 Kyona Labs plan withdrawn)

Horizon 06-Oct-2026 → 31-Mar-2027 · **Launch** Mon 19-Oct-2026 · **Gates** 18-Nov · 18-Dec · 31-Jan · 31-Mar

Built on four research passes in this chat (monetization & policy; tools & pipeline; topic calendar; faceless generative niches). **v2.2 change:** voice, dubbing and sound effects move from ElevenLabs direct to the Runway API (same ElevenLabs models, one credit pool, no subscription). Companion files: `kiln-prompts.md` (eight prompts, each labelled with its destination) and `kiln-topic-calendar.csv` (164 dated Shorts, 12 long-form episodes, 16 bank slots).

What changes from v1, and what we are building

v1 assumed real Kyona Labs footage. v2 drops Kyona entirely: no face, no studio, no products. Every frame is generated or rendered. That removes shooting time and the strongest originality signal, so **the IP has to carry it**. The channel is an original stickman workplace sitcom, **Bureau of Reality**: a recurring cast who run the department that keeps physics, time, history and myth working. Each Short explains one real mechanism through a character conflict.

Automation is organised around one idea: **MCP is the control plane, APIs are the data plane**. People and agents talk to the pipeline through one Kiln MCP server with scoped permissions; generation, voice, rendering and publishing run in Kiln workers against provider APIs, where cost, retries and limits are controlled.

THE BET IN ONE LINE

Own a cast of characters people recognise. Fully AI channels that survived 2025–26 had owned IP, a named human writer and episodes that genuinely differ; the ones terminated were template farms, borrowed-IP recaps, religious slop and kids content. Shorts build the cast; long-form, memberships and brand deals pay for it later.

Ten decisions

#	Decision	Choice	Why
1	Channel	One English channel, "Bureau of Reality", Brand Account, not made for kids	Termination risk is portfolio-wide; kids-coded AI is the most scrutinised category in 2026
2	Format	Stickman sitcom with 7 recurring characters; one real mechanism per episode; 45–70 s	Reuses your proven Higgsfield stickman style; owned IP; brand-safe edutainment
3	Series	Seven weekday series (Incident, Desk Tour, Pip's First Year, Archive, Myth Desk, Deep Desk, Complaint Box)	Variety is structural, not hoped for
4	Excluded	Finance/health hosts, kids style, devotional mythology, true crime, politics, borrowed IP	Each maps to a named 2026 enforcement pattern
5	Volume	1/day for 60 days; then 2/day on Tue/Thu/Sat only if Gate 2 passes	No official safe number; this keeps you far from farm patterns
6	Human input	You approve each premise and pick or rewrite the punchline (Tap 1), then approve the cut (Tap 2): ~10 min a day each	Authorship trail for appeals; the "human writer" signal
7	Control plane	Custom Kiln MCP server on Vercel, added to Claude as a custom connector; scoped tokens	One interface for you (phone), scheduled tasks, Routines and Claude Code
8	Data plane	Trigger.dev v4 + Supabase queue → Higgsfield, Gemini Veo, Runway (voice, dubs, Act-Two), Three.js render, YouTube/IG APIs	Per-second billing, refunds, retries, caps; MCP connectors bill web-plan credits
9	Languages	English first, auto-dub off; from day 30 add custom Hindi, Spanish and Portuguese tracks to the top 20%	Comedy timing breaks in unedited auto-dubs; text kept on a separate layer
10	Budget	Cap ₹15,000/month to Gate 2, then ₹25,000; per-Short cap ₹150	Expect under ₹1 lakh total income by March unless a series breaks out

Bureau of Reality

World: white chalk-line stick figures, one accent colour each, on deep navy blueprint paper. Mechanisms appear as "blueprint overlays" rendered by Kiln in Three.js — cheap, deterministic, re-renderable per language. Adult office-satire tone, never kid-coded.

Cast

Character	Role	Visual lock	Voice brief (preset audition)	Owns
Pip	New intern, audience surrogate	White line, cyan scarf, oversized lanyard	Bright, fast, curious	Pip's First Year (serialised)
Marlo	400-year Gravity Desk veteran	Amber line, floating coffee mug	Deadpan baritone	Incidents, Deep Desk
Mrs. Iyer	Head of Time & Calendars	Magenta line, reading glasses, filter-coffee tumbler	Crisp Indian-English, dry wit	Every festival and calendar episode
Nib	Archivist	Graphite line, pencil-shaped body	Whispery, conspiratorial	The Archive (no true crime)
Kaz	Myth Desk liaison	Lantern head; glow changes by culture	Warm storyteller	Myth Desk (comparative, respectful)
Director Ohm	Never-seen boss	A humming brass desk lamp	Filtered memo voice	Stakes, cold opens
Complaint Box	Reads viewer comments	Talking wooden suggestion box	Gravelly	Sunday viewer questions
The Auditor	Season 2 antagonist	Red line, clipboard	Clipped, corporate	The Audit arc (from 06-Jan)

VOICES ARE PRESETS, NOT DESIGNS

The Runway API speaks with Runway's preset voices on ElevenLabs models; its voice design and cloning work only for Runway avatars. Each character gets one preset, chosen by audition and locked in characters.json. **Audition gate:** if two characters are hard to tell apart, or a preset is obviously in use on other channels, move that character to ElevenLabs direct (Creator, \$22/month) behind the same voice driver. Name mentions per 1k views is the live check.

Week and episode shape

Day	Series	Lead	Template
Mon	Incident Report	Pip + Marlo	Accident → consequence ladder → real-world fact
Tue	Desk Tour	Rotating desk head	Tour → three-beat mechanism → twist
Wed	Pip's First Year	Pip	Serialised arc, one idea per episode; YouTube Shorts series when available to you
Thu	The Archive	Nib	Case file → evidence → best current explanation
Fri	Myth Desk	Kaz	Myth → parallel in another culture → meaning
Sat	Deep Desk	Marlo / Mrs. Iyer	Visual paradox → explanation (space, ocean)
Sun	Complaint Box	Complaint Box	Real viewer comment → answer; long-form every other Sunday

Beat sheet (45–70 s): 0–2 s visual paradox cold open · 2–8 s stakes · 8–40 s three mechanism beats with Three.js overlays · 40–55 s one sourced real-world fact (URL logged) · 55–60 s button gag and a loop line back to frame one.

Variation, enforced in code

- Each brief differs from every one of the last 14 episodes on at least **4 of 7 axes**: series, lead, desk, premise type, structure variant, ending type, music bed.
- Script embedding similarity against the last 60 episodes must be **< 0.85** (tunable); any hook archetype at most twice a week; no catchphrase more than once a week.
- Policy lint rejects finance or health advice, politics, real living people, franchises and brands, true crime, kid-coded styling.
- Every episode stores your premise approval and punchline choice with timestamps — the authorship log for any appeal.

Languages

One channel, English master. Auto-dubbing off at launch. Captions and on-screen text are rendered as a separate layer so every Short has a clean master. From day 30 the weekly Routine queues the top 20% of Shorts for ElevenLabs dubbing through the Runway API in Hindi, Spanish and Portuguese (Brazil), each character keeping a matched voice; you add the tracks as multi-language audio in Studio. Hindi has the lowest competition; Spanish and Portuguese pay better.

03 · AUTOMATION AND MCP

MCP is the control plane, APIs are the data plane

Every actor — you on your phone, a Claude scheduled task, a Claude Code Routine, or Claude Code while building — reads and changes the pipeline through one Kiln MCP server. That server never generates video itself; it writes to Supabase and starts Trigger.dev tasks (Kiln's existing orchestrator). Workers call provider APIs with per-job cost estimates, caps and retries. This separation is what makes "fully automated" safe: agents can propose and run work, but only your token can approve or publish.




FIGURE 1 · CONTROL PLANE AND DATA PLANE. ONLY THE PURPLE ACTOR HOLDS APPROVE AND PUBLISH SCOPES.

MCP servers and what each one is for

MCP server	Connected in	Job	In production?
Kiln MCP (you build it — Prompt A)	Claude (custom connector), Claude Code, Routines	Single control plane: calendar, briefs, approvals, regenerations, cut review, publish bundles, metrics, costs, caps, kill switch	Yes — the only way agents touch the pipeline
Supabase MCP	Claude Code	Migrations, schema diffs, logs, advisors, ad-hoc analytics SQL	Build and ops only
Higgsfield MCP	Claude chat	Character sheets, style exploration, preset tests — with stop-and-ask before credits	No — bills web-plan credits
Slack MCP	Claude scheduled tasks	Morning desk digest, cuts-ready ping, cost and policy alerts	Yes, notify only
Google Calendar MCP	Claude scheduled tasks	Biweekly long-form review blocks	Optional
Runway MCP (official, mcp.runwayml.com/mcp , add as a custom connector) + Runway app	Claude chat	Image and video exploration and Act-Two trials on your app credits ; audition preset voices in the Runway app. The MCP exposes no voice tools	No — no retries or cost rows; production uses the Runway API (Prompt H)
Notion	—	Not used. Supabase is the single source of truth	—

Kiln MCP — tool surface

Tool	Does	Scope
calendar_upcoming(days)	Slots, series, topics, seasonal tags, bank	read
briefs_create_batch(briefs[])	Writes drafts as needs_approval; runs variation + policy checks inline	agent
brief_get(id) · briefs_pending()	Premise, 3 punchlines, beat sheet, cost estimate, check results	read

Tool	Does	Scope
brief_approve(id, punchline, edits)	Tap 1. Logs your text and timestamp; starts the episode workflow	approver
brief_reject(id, reason)	Kills or returns for rewrite	approver
variation_check(brief) · policy_lint(script)	Dry-run checks for any draft	read
episode_status(id)	Stage, jobs, QC scores, spend so far	read
shot_regenerate(episode, shot, note)	Re-queues one shot within the cap	agent
cut_approve(id) · cut_reject(id, note)	Tap 2	approver
publish_bundles() · mark_scheduled(id, at)	Ready-to-schedule MP4 + metadata; records Studio scheduling	approver
metrics_summary(range) · top_performers(n)	Views, viewed-vs-swiped, APV, subs/1k, name mentions	read
complaint_candidates()	Best viewer questions for Sunday episodes	read
dub_queue(add/list)	Top-20% Shorts for hi / es / pt-BR tracks	agent
costs_ledger(range) · caps_set(...)	Spend by stage and provider; per-Short, daily and monthly caps	read / approver
kill_switch(on)	Stops all generation and publishing	approver

Two tokens: **approver** (yours, used from Claude chat and the control room) and **agent** (Routines and scheduled tasks). An agent token can never approve its own work or publish — enforced server-side, not by prompt.

The daily loop

IST	Runs on	What happens	You
06:00	Routine "Showrunner" (Claude Code, cloud)	Reads calendar D+3, yesterday's metrics and the variation ledger through Kiln MCP; drafts briefs with 3 punchlines each; regenerates up to twice on failed checks	—
07:30	Scheduled task "Morning desk" (Claude app)	Pulls briefs_pending, posts them to your Slack DM with yesterday's numbers	—
any time	You in Claude chat (phone)	"Approve 1 with B; rewrite 2's punchline to ..." → brief_approve	~10 min
on approve	Trigger.dev task	Script polish → shotlist → router → Higgsfield/Veo jobs from locked character refs → Three.js overlays → Runway voices + word alignment → QC → master + caption layer	—
20:30	Slack ping	Cuts ready; review in the control room or ask Claude to show and approve	~10 min
slot time	Kiln publisher	Instagram Reels posts automatically. YouTube: private upload + bundle until the API audit passes, then scheduled public	Mon/Thu 15 min until audit
+1h/24h/72h/7d	Workflow sleeps	Metric pulls into Supabase; bandit updates hook and length arms	—
Mon 06:30	Routine "Weekly review"	Gates, cost vs cap, sequel picks into the bank, dub queue, date re-verification, strategy memo	read 15 min
Thu (biweekly)	Routine "Long-form"	Assembles the Sunday episode plan from aired Shorts + new connective scenes → one brief for approval	approve

Generation routing for this channel

Shot	Engine	≈ \$/s	Rule
Blueprint overlays, diagrams, text	Kiln Three.js render	0.00	≥ 50% of runtime
Character acting beats	Kling 3.0 Std image-to-video from locked character frames (Higgsfield API)	0.084	≤ 8 s per Short
Your acted performance (optional)	Runway Act-Two	0.05	Strong authorship signal
Money shot (ocean, space, fireworks)	Veo 3.1 Lite / Fast (Gemini API)	0.05–0.10	Max 1 per Short
Voices	Runway API, Eleven v3 text-to-speech, one locked preset per character; word timings by forced alignment on the worker (Runway returns none)	–	≈ \$0.18 per Short; preset IDs locked
Dubs, sound effects	Runway API (ElevenLabs Dubbing, Text to Sound Effects)	–	Dubbing rate to verify before day 30

04 • ECONOMICS

Costs, money and the honest outlook

Per Short	USD	₹	Basis
Character shots	1.01	89	8 s Kling Std × 1.5 reroll factor
LLM (brief, polish, judge, lint) + QC	0.18	16	Sonnet bulk, Opus judge, Haiku checks
Voice	0.18	16	~900 characters × 1 credit / 50 chars (Eleven v3), \$0.01 a credit
Optional money shot	0.30	26	4 s Veo Lite × 1.5
Typical range	1.37–1.67	121–147	Cap ₹150 — only ₹3 headroom with a money shot, so the router drops it to an overlay when a reroll would cross the cap

Monthly	Phase 1–2 • 1/day	Phase 3+ • ~43/mo + dubs	Note
Shorts (₹140 average, voice included)	₹4,200	₹6,020	30 vs 43 Shorts
Long-form, 2 a month	₹3,000	₹3,000	New connective scenes only
Runway API — dubs, Act-Two	₹1,000	₹2,500	Budget, not a quote: Runway does not publish a dubbing rate — measure it on the first dub
Vercel Pro + Supabase Pro	₹3,960	₹3,960	App, MCP endpoint, storage; Trigger.dev usage on top
Total, excluding Claude	₹12,160	₹15,480	Caps ₹15k and ₹25k; ₹88/USD; ~₹5.5k/month below v2.1 at Phase 3

One-off: Higgsfield API wallet \$25 (unlocks ~10 concurrent jobs; funds last a year). Use your existing Higgsfield web credits for character sheets. Runway API credits \$25 on the developer portal (~130 Shorts of voice plus Act-Two trials). **Runway app credits do not transfer to the API** — spend them on the voice audition in the Runway app and on exploration through the Runway MCP.

Where money comes from

Stream	Unlocks	Realistic by Mar-2027
Patreon / memberships	Patreon from day 1 (scripts, character sheets, early episodes, complaint voting); memberships at the expanded tier	₹0–25k/month

Stream	Unlocks	Realistic by Mar-2027
Expanded YPP tier	500 subs + 3M Shorts views/90d — unchanged in 2027. Target January	Small, but unlocks fan funding
Shorts ad pool	≥ 10M qualified views/90d from 01-Feb-2027	~\$800–3,300 a quarter if reached
Long-form ads	Old bar to 31-Jan (4,000 hours) is unlikely; plan for 8,000 hours	Likely ₹0 in window
Affiliates	Day 1 — books and education apps tied to each episode	Tens of dollars
Brand deals, character licensing	Once the cast is recognisable (judgement: 50k engaged subs)	Likely ₹0 in window

HONEST OUTLOOK

Assume less than ₹1 lakh of total income by 31-Mar-2027 unless a series breaks out, against roughly ₹85k of spend (excluding Claude) over the same period. Treat the six months as building an IP asset with a measured audience. The plan is designed so that this conservative case costs no more than the caps.

Build, bank, launch, gate

The research calendar launched on 12-Oct with 30 pre-made Shorts. That isn't realistic while Kiln is still being extended, so v2 launches **Mon 19-Oct** with **14 Shorts banked**. The six evergreen 12-18 Oct topics and the leap-year episode move into the bank as cover for QC failures; Pip's orientation becomes S1E1 on Wed 21-Oct. All seasonal dates are unchanged.

Phase	Dates	Scope	Exit gate
0 • Build & bank	06 → 18-Oct	Accounts, API audit filed, Kiln retarget + Kiln MCP (Prompt A), voice switch (Prompt H), cast locked (Prompt B), preset voices auditioned and locked, 14 Shorts + pilot long-form banked	14 approved cuts; both taps work from Claude chat; cost ≤ ₹150 each
1 • Launch	19-Oct → 17-Nov	1/day; Halloween, DST, Diwali run; Routines and morning desk live; IG mirror on	Day-30 review 18-Nov: replace weakest series slot; start dubs on top 20%
2 • Prove	18-Nov → 17-Dec	1/day + dubs; Archive Vol. 1, mid-season cut, Christmas run	Gate 18-Dec: 2/day on Tue/Thu/Sat only if KPIs hold
3 • Season 2	18-Dec → 31-Jan	S1 finale, The Audit arc, Sankranti/Pongal, expanded-tier push, Patreon tiers live	Gate 31-Jan: KPIs held at volume; spend inside cap
4 • Compound	01-Feb → 31-Mar	Valentine's, Holi, S2 finale, S3 start; first brand outreach if recognisable	Gate 31-Mar: continue, or open channel 2 ("Mira Out of Time") only if channel 1 is stable

Phase 0, day by day

When	You	Claude
Tue 06-Oct	Brand Account channel "Bureau of Reality"; Google Cloud project + file YouTube API audit ; Higgsfield API \$25; Gemini billing; Runway API credits \$25	Prompt A starts in Claude Code (one long run)
Wed 07-Oct	Prompt B in Claude chat with Higgsfield: approve each sheet before credits are spent	Sprints 0-2: retarget, schema, Kiln MCP
Thu-Fri 08-09	Lock reference frames; audition and lock 8 preset voices; write premises + punchlines for the first 14	Sprints 3-4: workflows, router; Prompt H once the build lands
Sat-Sun 10-11	Add Kiln MCP as a custom connector in Claude; first dry run from your phone	Sprints 5-6: QC, publish, metrics
Mon-Fri 12-16	Approve and review 2-3 Shorts a day	Sprint 7-8: control room, dubs; Routines C, E, F and task D created
Sat 17-Oct	Approve the pilot long-form (airs 25-Oct)	Bank check: 14 cuts
Sun 18-Oct	Gate 0 checklist; Meta app for Instagram	—
Mon 19-Oct	Launch — "What if the Moon vanished"	

FALLBACK

If Kiln isn't producing reliably by 14-Oct, make the launch 14 with Kiln's existing script, shotlist and composition stages plus direct Higgsfield API calls, and keep the build running. Do not slip launch past 26-Oct: Halloween and Diwali episodes are date-locked.

Your week, KPIs and gates

Activity	When	Hours / week
Tap 1 — premises and punchlines, from Claude chat	Daily, 10 min	1.2
Tap 2 — final cuts	Daily, 10 min	1.2
Comments and Complaint Box picks	Daily, 10 min	1.2
Long-form review (biweekly)	Alternate Saturdays	0.8
Studio scheduling until the API audit passes	Mon + Thu	0.5
Weekly memo + experiments	Mon	0.5
Total		5.4

KPI	Gate	Measured	If missed
Viewed vs swiped away	≥ 70%	Median, last 20 Shorts	Rework cold opens; kill formats under 60% after 10 tries
Average percentage viewed	≥ 70%	Median, 45–70 s	Tighten beats 2–3; shorter default length
Subscribers per 1k views	≥ 1.0	Rolling 28 days	More serialised Pip episodes
Character-name mentions per 1k views	rising	Comment mining	The IP isn't landing; sharpen character conflict
Cost per published Short	≤ ₹150	Cost ledger	More Three.js runtime; drop money shots
Policy warnings	0	Studio + policy log	Kill switch; audit last 20

KILL RULE

Viewed-vs-swiped under 60% after 40 Shorts, or any monetization or "generic or repetitive" warning: pause a week, change format, never answer it with more volume.

Risks

Risk	Likely	Mitigation
"Generic or repetitive" demonetization of a fully generated channel	Medium	Owned cast, 4-of-7 variation rule, similarity block, your premise and punchline on every episode, authorship log
An agent publishes or overspends	Low	Approve, publish and caps are approver-scope only; server-side caps; kill switch
Character drift across shots	High early	Image-to-video from locked frames only; QC similarity vs reference; reroll ≤ 2
Preset voices make the cast generic	Medium	Audition gate; any character can move to ElevenLabs direct behind the same driver; watch name mentions
Kid-coded perception (stickman)	Medium	Adult office satire, not-made-for-kids setting, no nursery colours or voices
Myth episodes offend	Medium	Comparative and interpretive framing, "said to mean"; no devotional content
Vendor shutdown or price change (Sora shut 24-Sep-2026)	Certain over time	Every engine behind a swappable adapter; fal.ai failover; costs logged per job
API audit delay	High	Bundles + Studio scheduling; audit is an upgrade, not a dependency

Risk	Likely	Mitigation
Faceless reach penalty (creator claim, unconfirmed)	Unknown	Recognisable characters stand in for a face; measure name mentions
Your time across other projects	High	14-Short bank; taps work from phone; pause beats slop

This week

1. **Today:** create the Brand Account channel, file the YouTube API audit, fund the Higgsfield API wallet (\$25), enable Gemini billing, buy Runway API credits (\$25, developer portal).
2. **Today:** copy `kiln-topic-calendar.csv` into the Kiln repo at `data/` and paste **Prompt A** into Claude Code.
3. **Wed:** run **Prompt B** in Claude chat with the Higgsfield connector to lock the cast.
4. **Thu–Fri:** run **Prompt H** in Claude Code after the build lands; audition and lock the eight preset voices; write the first 14 premises and punchlines.
5. **Weekend:** add the Kiln MCP server as a custom connector in Claude; approve your first Short from your phone.
6. **Mon 12-Oct:** create Routines C, E, F and scheduled task D; open Prompt G in Claude Design.

Sources

YouTube Blog — YPP 2027 changes; Made on YouTube (23-Sep-2026); auto-dubbing (Feb-2026). YouTube Help — Shorts monetization policies, channel monetization policies, multi-language audio. TeamYouTube Shorts update (01-Oct-2026). TechCrunch — Shorts series (23-Sep-2026). Kapwing AI Slop Report; The Hollywood Reporter (13-Jun-2026); The Next Web; Fortune (01-Apr-2026); Tubefilter; Variety (Sora, ElevenLabs × Hasbro); XDA Developers; Deadline via NetInfluencer; Higgsfield API help centre and blog (Sep-2026); Runway API usage tiers, pricing, changelog and "How do credits work" (Oct-2026); Gemini API pricing; Trigger.dev v4 docs; AIR Media-Tech and Fluxnote RPM data. Full citations are in the four research reports in this chat. Channel examples (Neural Viz, Chloe VS History) are point-in-time figures.